

A* GPGPU Implementations for Path-Finding

Florian Klemme

July 26, 2017

Overview

The A* search algorithm

- The basics

- GPGPU approaches

Priority queue as binary heap

- Dispatch memory access

- Binary heap functions

- Push, pop and find

Multi-agent A*

- Kernel snippets

Parallel GA*

- Data flow

- Kernel selection

- Use of synchronization mechanisms

Conclusion

- Lessons learned

- Demo performance

- Boost Compute

The A* search algorithm

- ▶ **A*** is a *best-first search* algorithm
- ▶ Chooses minimal node $f(n) = g(n) + h(n)$
- ▶ Needs *priority queue* to select next node
- ▶ Needs look-up data structure for visited nodes

Animation CC BY 3.0
Subh83 on wikimedia.org

Multi-agent A* vs. parallel GA*

Problem: It's hard to parallelize the priority queue!

Multi-agent A* [1]

- ▶ Per-agent local priority queue
- ▶ Node and edge information can be shared. Open and closed list has to be stored for each agent.
- ▶ **Local memory is a scarce resource!** I've seen 16, 32 and 48 kilobyte.

Parallel GA* [2]

- ▶ Multiple independent priority queues
- ▶ Expand multiple nodes at the same time
- ▶ The challenge is **handling of duplicate nodes!**
- ▶ Again, bound by memory

Priority queue: Dispatch memory access

```
typedef struct {
    __local   uint_float *localMem;
    const     size_t      localSize;
    __global  uint_float *globalExt;
             size_t      size;
} OpenList;

uint_float _read_heap(OpenList *open, size_t index) {
    return index < open->localSize ?
        open->localMem[index] :
        open->globalExt[index - open->localSize];
}

void _write_heap(OpenList *open, size_t      index,
                 uint_float value) {
    if (index < open->localSize)
        open->localMem[index] = value;
    else
        open->globalExt[index - open->localSize] = value;
}
```

Priority queue as binary heap

We need the common...

- ▶ push
- ▶ top & pop

But also...

- ▶ find
- ▶ update

```
void push(OpenList *open, uint value, float cost) {  
    _push_impl(open, &open->size, value, cost);  
}
```

```
void update(OpenList *open, size_t index, uint value,  
            float cost) {  
    _push_impl(open, &index, value, cost);  
}
```

Binary heap: Push

```
void _push_impl(OpenList *open, size_t *size,
                uint      value, float   cost) {
    size_t index = (*size)++;

    while (index > 0) {
        size_t parent = (index - 1) / 2;

        uint_float pValue = _read_heap(open, parent);
        if (cost < pValue.second) {
            _write_heap(open, index, pValue);
            index = parent;
        } else
            break;
    }

    _write_heap(open, index, (uint_float){value, cost});
}
```

Binary heap: Pop

```
void pop(OpenList *open) {
    uint_float value = _read_heap(open, --(open->size));
    size_t      index = 0;

    while (index < open->size / 2) {
        size_t      child = index * 2 + 1;
        uint_float cValue = _read_heap(open, child);
        if (child + 1 < open->size) {
            uint_float c1Value = _read_heap(open,
                                             child + 1);

            if (c1Value.second < cValue.second) {
                ++child;
                cValue = c1Value;
            }
        }

        if (cValue.second < value.second) {
            _write_heap(open, index, cValue);
            index = child;
        } else break;
    }
}
```


Binary heap: Find

This could certainly be improved! Now it's basically a simple breadth-first search.

```
uint find(OpenList *open, uint value) {  
    for (uint index = 0; index < open->size; ++index) {  
        uint_float iValue = _read_heap(open, index);  
        if (iValue.first == value)  
            return index;  
    }  
  
    return open->size;  
}
```

Alternative: Except duplicate entries and ignore them later on.
Would cost extra memory!

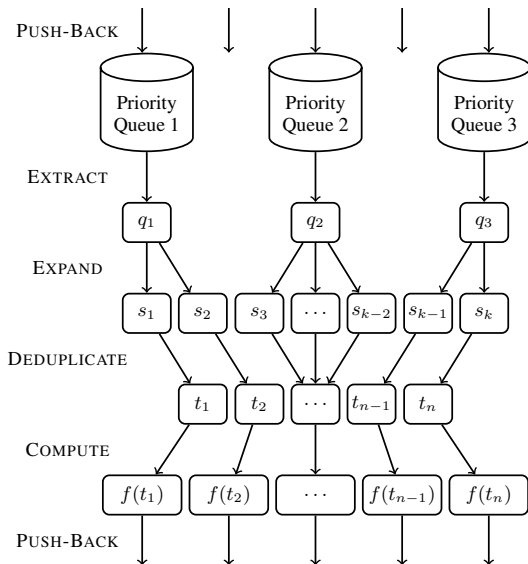
Multi-agent A* kernel

```
push(&open, source, 0.0f);
while (open.size > 0) {
    const uint current = top(&open); pop(&open);
    /* ... */
    const uint2 range = adjacencyMap[current];
    for (uint edge = range.x; edge != range.y; ++edge) {
        /* ... */
        const uint nbIndex = find(&open, nbNode);

        if (nbIndex < open.size &&
            nbInfo.totalCost <= nbTotalCost) continue;
        /* ... */

        if (nbIndex < open.size)
            update(&open, nbIndex, nbNode, nbTotalCost +
                nbHeuristic);
        else
            push(&open, nbNode, nbTotalCost + nbHeuristic);
    }
}
```

Data flow in parallel GA*



Run kernels, in order:

- ▶ clearSList
- ▶ extractAndExpand
- ▶ clearTList
- ▶ duplicateDetection
- ▶ exclusive_scan
- ▶ compactTList
- ▶ computeAndPushBack

Image from [2].

GA* Pseudocode

Algorithm 1 GA*: Parallel A* search on a GPU

```
1: procedure GA*( $s, t, k$ )  
     $\triangleright$  find the shortest path from  $s$  to  $t$  with  $k$  queues  
2:   Let  $\{Q_i\}_{i=1}^k$  be the priority queues of the open list  
3:   Let  $H$  be the closed list  
4:   PUSH( $Q_1, s$ )  
5:    $m \leftarrow \text{nil}$   $\triangleright m$  stores the best target state  
6:   while  $Q$  is not empty do  
7:     Let  $S$  be an empty list  
8:     for  $i \leftarrow 1$  to  $k$  in parallel do  
9:       if  $Q_i$  is empty then  
10:        continue  
11:       end if  
12:        $q_i \leftarrow \text{EXTRACT}(Q_i)$   
13:       if  $q_i.\text{node} = t$  then  
14:         if  $m = \text{nil}$  or  $f(q_i) < f(m)$  then  
15:            $m \leftarrow q_i$   
16:         end if  
17:         continue  
18:       end if  
19:        $S \leftarrow S + \text{EXPAND}(q_i)$   
20:     end for  
21:     if  $m \neq \text{nil}$  and  $f(m) \leq \min_{q \in Q} f(q)$  then  
22:       return the path generated from  $m$   
23:     end if  
24:      $T \leftarrow S$   
25:     for  $s' \in S$  in parallel do  
26:       if  $s'.\text{node} \in H$  and  $H[s'.\text{node}].g < s'.g$  then  
27:         remove  $s'$  from  $T$   
28:       end if  
29:     end for  
30:     for  $t' \in T$  in parallel do  
31:        $t'.f \leftarrow f(t')$   
32:       Push  $t'$  to one of priority queues  
33:        $H[t'.\text{node}] \leftarrow t'$   
34:     end for  
35:   end while  
36: end procedure
```

Run kernels, in order:

- ▶ clearSList
- ▶ extractAndExpand
- ▶ clearTList
- ▶ duplicateDetection
- ▶ exclusive_scan
- ▶ compactTList
- ▶ computeAndPushBack

Pseudocode from [2].

Use of synchronization mechanisms: return code

```
__kernel void extractAndExpand(/* ... */  
                               __global uint *returnCode) {  
    /* ... */  
    if (openSize == 0) {  
        atomic_min(returnCode, 2);  
        return; // failure: no path found!  
    }  
  
    /* ... */  
    if (current == destination) {  
        atomic_min(returnCode, 0);  
        return; // success: path found!  
    }  
  
    /* ... */  
    atomic_min(returnCode, 1); // still running...  
}
```

Use of synchronization mechanisms: inserting elements

```
--kernel void duplicateDetection(/* ... */) {  
    /* ... */  
    if (nodeInfo.closed == 1 &&  
        nodeInfo.totalCost < current.totalCost)  
        return; // better candidate already in open list  
  
    const uint hash = current.node % hashTableSize;  
    const uint old = atomic_xchg(hashTable + hash,  
                                   current.node);  
  
    if (old == current.node)  
        return; // node has already been added  
  
    __global Info *tlist =  
        tlistChunks + GID.x * slistChunkSize;  
    const uint index = atomic_inc(tlistSizes + GID.x);  
    tlist[index] = current;  
}
```

Updating information: Transaction approach

```
#pragma OPENCL EXTENSION cl_khr_int64_base_atomics : enable
```

```
typedef struct {  
    uint   closed;    uint   node;  
    float  totalCost; uint   predecessor;  
} Info;
```

```
// Update totalCost and pred. as one 64 bit transaction.  
ulong *curCostPred = (ulong *) &current.totalCost;  
__global ulong *infoCostPred =  
    (__global ulong *) &info[current.node].totalCost;  
ulong oldCostPred = atom_xchg(infoCostPred, *curCostPred);  
float *oldCost = (float *) &oldCostPred;  
  
// assert: current.totalCost > 0.0f  
while (*oldCost != 0.0f && *oldCost < current.totalCost) {  
    // The old entry was better. Swap back!  
    *curCostPred = oldCostPred;  
    oldCostPred = atom_xchg(infoCostPred, *curCostPred);  
}
```

Lessons learned

- ▶ OpenCL is hard to make portable and reliable
 - ▶ Compiler crashes
 - ▶ Wrong work sizes (even through API)
- ▶ Performance depends on actual hardware
- ▶ CPU A* is hard to beat, random obstacles have a huge impact

Multi-agent A*

- ▶ Work items $>$ local memory
- ▶ Make use of `__private` memory?

Parallel GA*

- ▶ Needs lots of memory, highly dependent on random obstacles
- ▶ Still lots of room for improvement!

Demo configuration on my secondary computer

```
OpenCL device: GeForce GT 520
----- CPU reference run...
CPU time for 2500 runs: 0.545225 seconds
----- GPU A* run...
GPU time for 2500 runs:
- Upload time: 0.0184138 seconds
- Kernel runtime: 1.05191 seconds
- Download time: 0.00187089 seconds
----- CPU reference run...
CPU time for graph (500, 500): 0.655936 seconds
----- GPU GA* run...
GPU time for graph (500, 500):
- Upload time: 0.00613427 seconds
- Kernel runtimes:
  - CompactTList: 0.0401437 seconds
  - ComputeAndPushBack: 0.129521 seconds
  - DuplicateDetection: 0.102979 seconds
  - ExtractAndExpand: 0.358814 seconds
  - compute::exclusive_scan: 0.505029 seconds
- Download time: 0.000844134 seconds
GPU GA*: Gold test failed!
- Path length CPU: 580, GPU: 581
- Path cost CPU: 728.528, GPU: 742.369
```

Boost Compute

Pure Template library, brings RAII, type-safety, common interface

```
// get the default compute device
compute::device gpu = compute::system::default_device();

// create a compute context and command queue
compute::context ctx(gpu);
compute::command_queue queue(ctx, gpu);

// generate random numbers on the host
std::vector<float> host_vector(1000000);
std::generate(host_vector.begin(), host_vector.end(), rand);

// create vector on the device
compute::vector<float> device_vector(1000000, ctx);

// copy data to the device
compute::copy(
    host_vector.begin(), host_vector.end(), device_vector.begin(), queue
);

// sort data on the device
compute::sort(
    device_vector.begin(), device_vector.end(), queue
);

// copy data back to the host
compute::copy(
    device_vector.begin(), device_vector.end(), host_vector.begin(), queue
);
```

<https://github.com/boostorg/compute>

References

- [1] Andre Silva, Fernando Rocha, Artur Santos, Geber Ramalho, and Veronica Teichrieb. Gpu pathfinding optimization. In *Games and Digital Entertainment (SBGAMES), 2011 Brazilian Symposium on*, pages 158–163. IEEE, 2011.
- [2] Yichao Zhou and Jianyang Zeng. Massively parallel a* search on a gpu. In *AAAI*, pages 1248–1255, 2015.

In case you want to check it out:

<https://github.com/Kruecke/ocl-astar>